

Azure Meetup

Kubernetes Hands-On!



VMWARE {code}



nirmata

Do this now!

- Signup for Azure

<https://azure.microsoft.com/en-us/free/>

- Signup for Nirmata

<https://try.nirmata.io>




Agenda

- 6:30 – 6:45: Kubernetes Cluster Concepts
- *6:45 – 7:45: Setup your Azure K8s cluster*
- 7:45 – 8:00: Kubernetes Application Concepts
- *8:00 – 9:00: Manage K8s applications with Nirmata*

Why Kubernetes?

- Containers and Microservices enable DevOps
 - Agility, Cost-savings, Portability
- Kubernetes has become the de-facto standard for container orchestration
 - Community-driven
 - Robust, scalable, and extensible




Kubernetes Cluster Concepts



Kubernetes Cluster

- Nodes
- Components
- Add-ons
- Cloud Provider
 - Networking
 - Storage



Master Nodes

- Master nodes run Kubernetes components
 - **kube-apiserver**: *front-end for the Kubernetes control-plane*
 - **etcd**: *datastore for the cluster*
 - **kube-controller-manager**: *controllers for routine cluster tasks*
 - **cloud-controller-manager**: *controllers specific to cloud providers*
 - **kube-scheduler**: *assigns Pods to nodes*



Each worker node runs

- Required:
 - **kubelet**: manages pods, executes liveness probes, reports pod and node status.
 - **kube-proxy**: network proxy; performs connection forwarding
 - **docker / rkt / containerd**: the container engine
- Optional
 - **add-ons**
 - **supervisord**
 - **fluentd**
- Your application pods



Common Add-ons

- **DNS**: serves DNS for Kubernetes components and containers. Consider as required.
- **Heapster**: provides container resource monitoring. Is used for Horizontal Pod Autoscaling.
- **Web UI**: dashboard to monitor and manage the cluster.

Networking

- K8s networking follows these principles:
 - All containers can communicate with all other containers without NAT
 - All nodes can communicate with all containers (and vice-versa) without NAT
 - The IP that a container sees itself as is the same IP that others see it as
- Each pod gets its own IP address
- CNI is the plugin model used by the Kubelet to invoke the networking implementation
- CNI plugins: Calico, Contiv, Flannel , GCE, ...



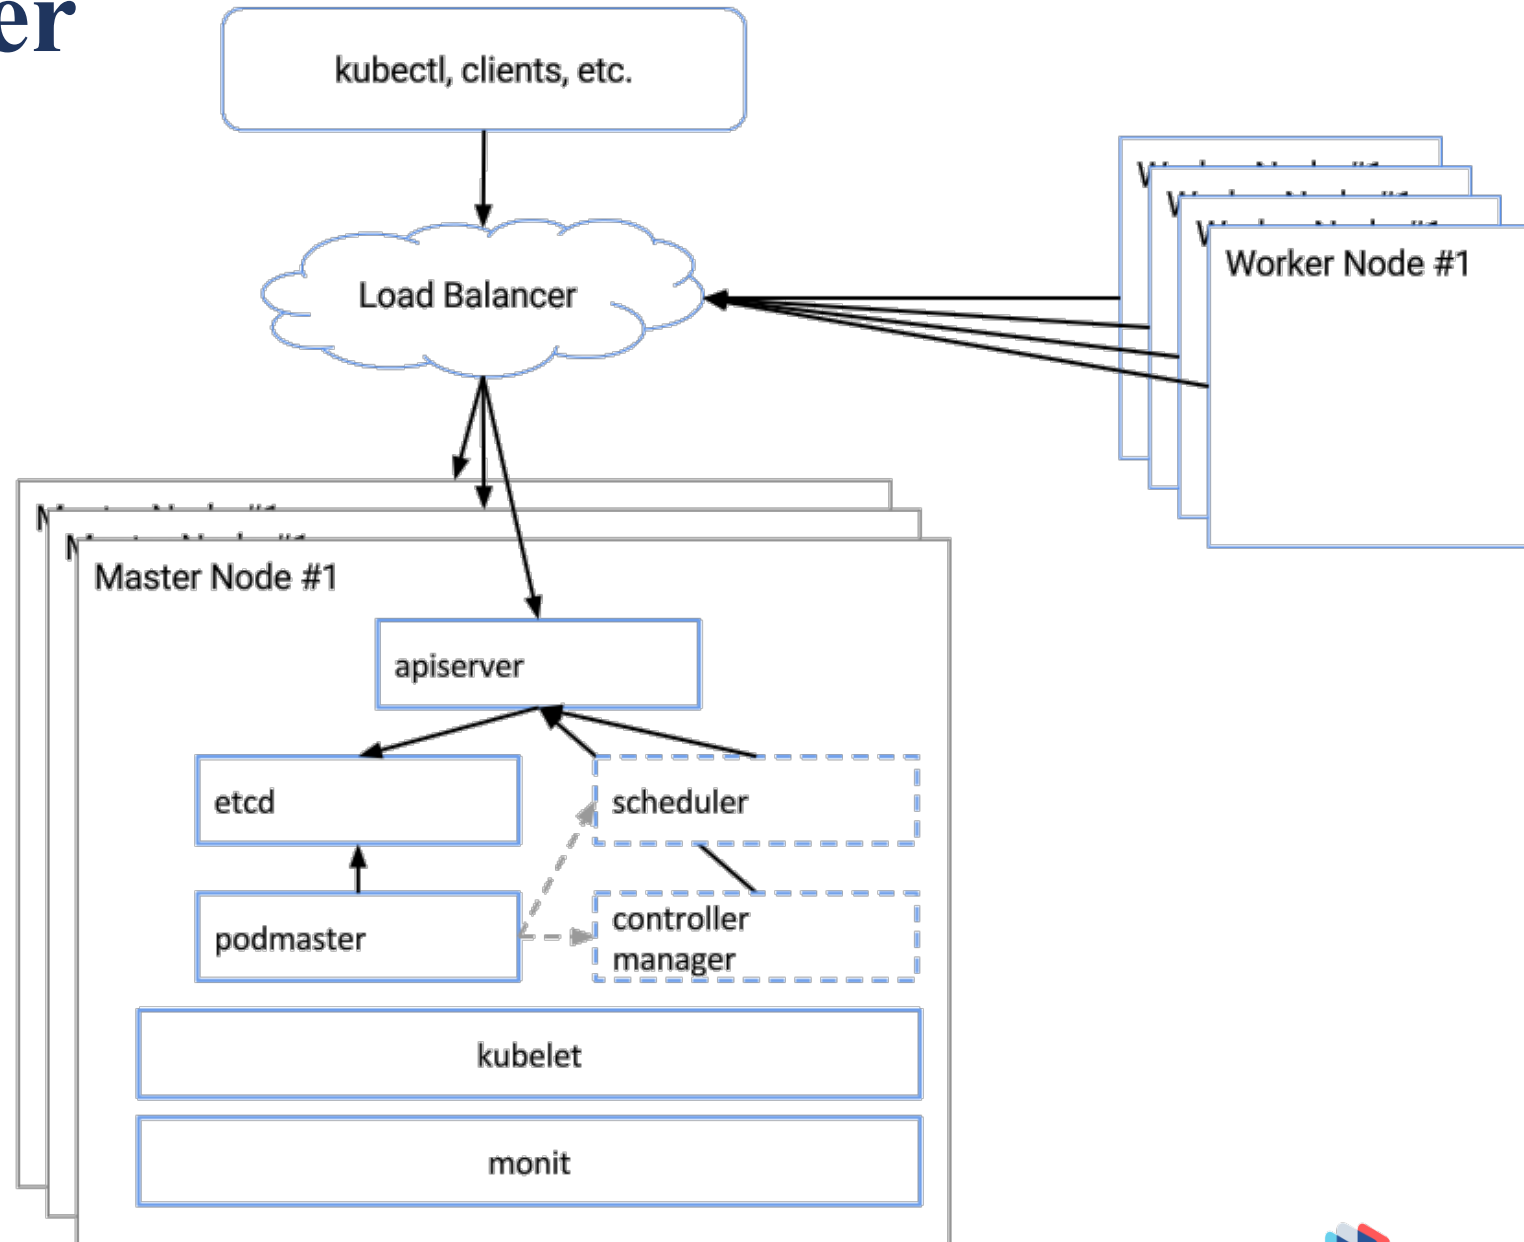
Storage

- Pods can contain one or more **Volumes**
 - *Volume types: emptyDir, hostPath, persistentVolumeClaim, secret, awsElasticBlockStore, AzureDiskVolume, ...*
- A **PersistentVolumeClaim** *requests a PersistentVolume that may be dynamically provisioned.*
 - Admins can configure **StorageClasses** for persistent volume claims like “bronze”, “silver”, or “gold”. A storage class has a **Provisioner**, like AzureDisk.



Kubernetes Cluster

- Nodes
- Components
- Add-ons
- Cloud Provider
 - Networking
 - Storage



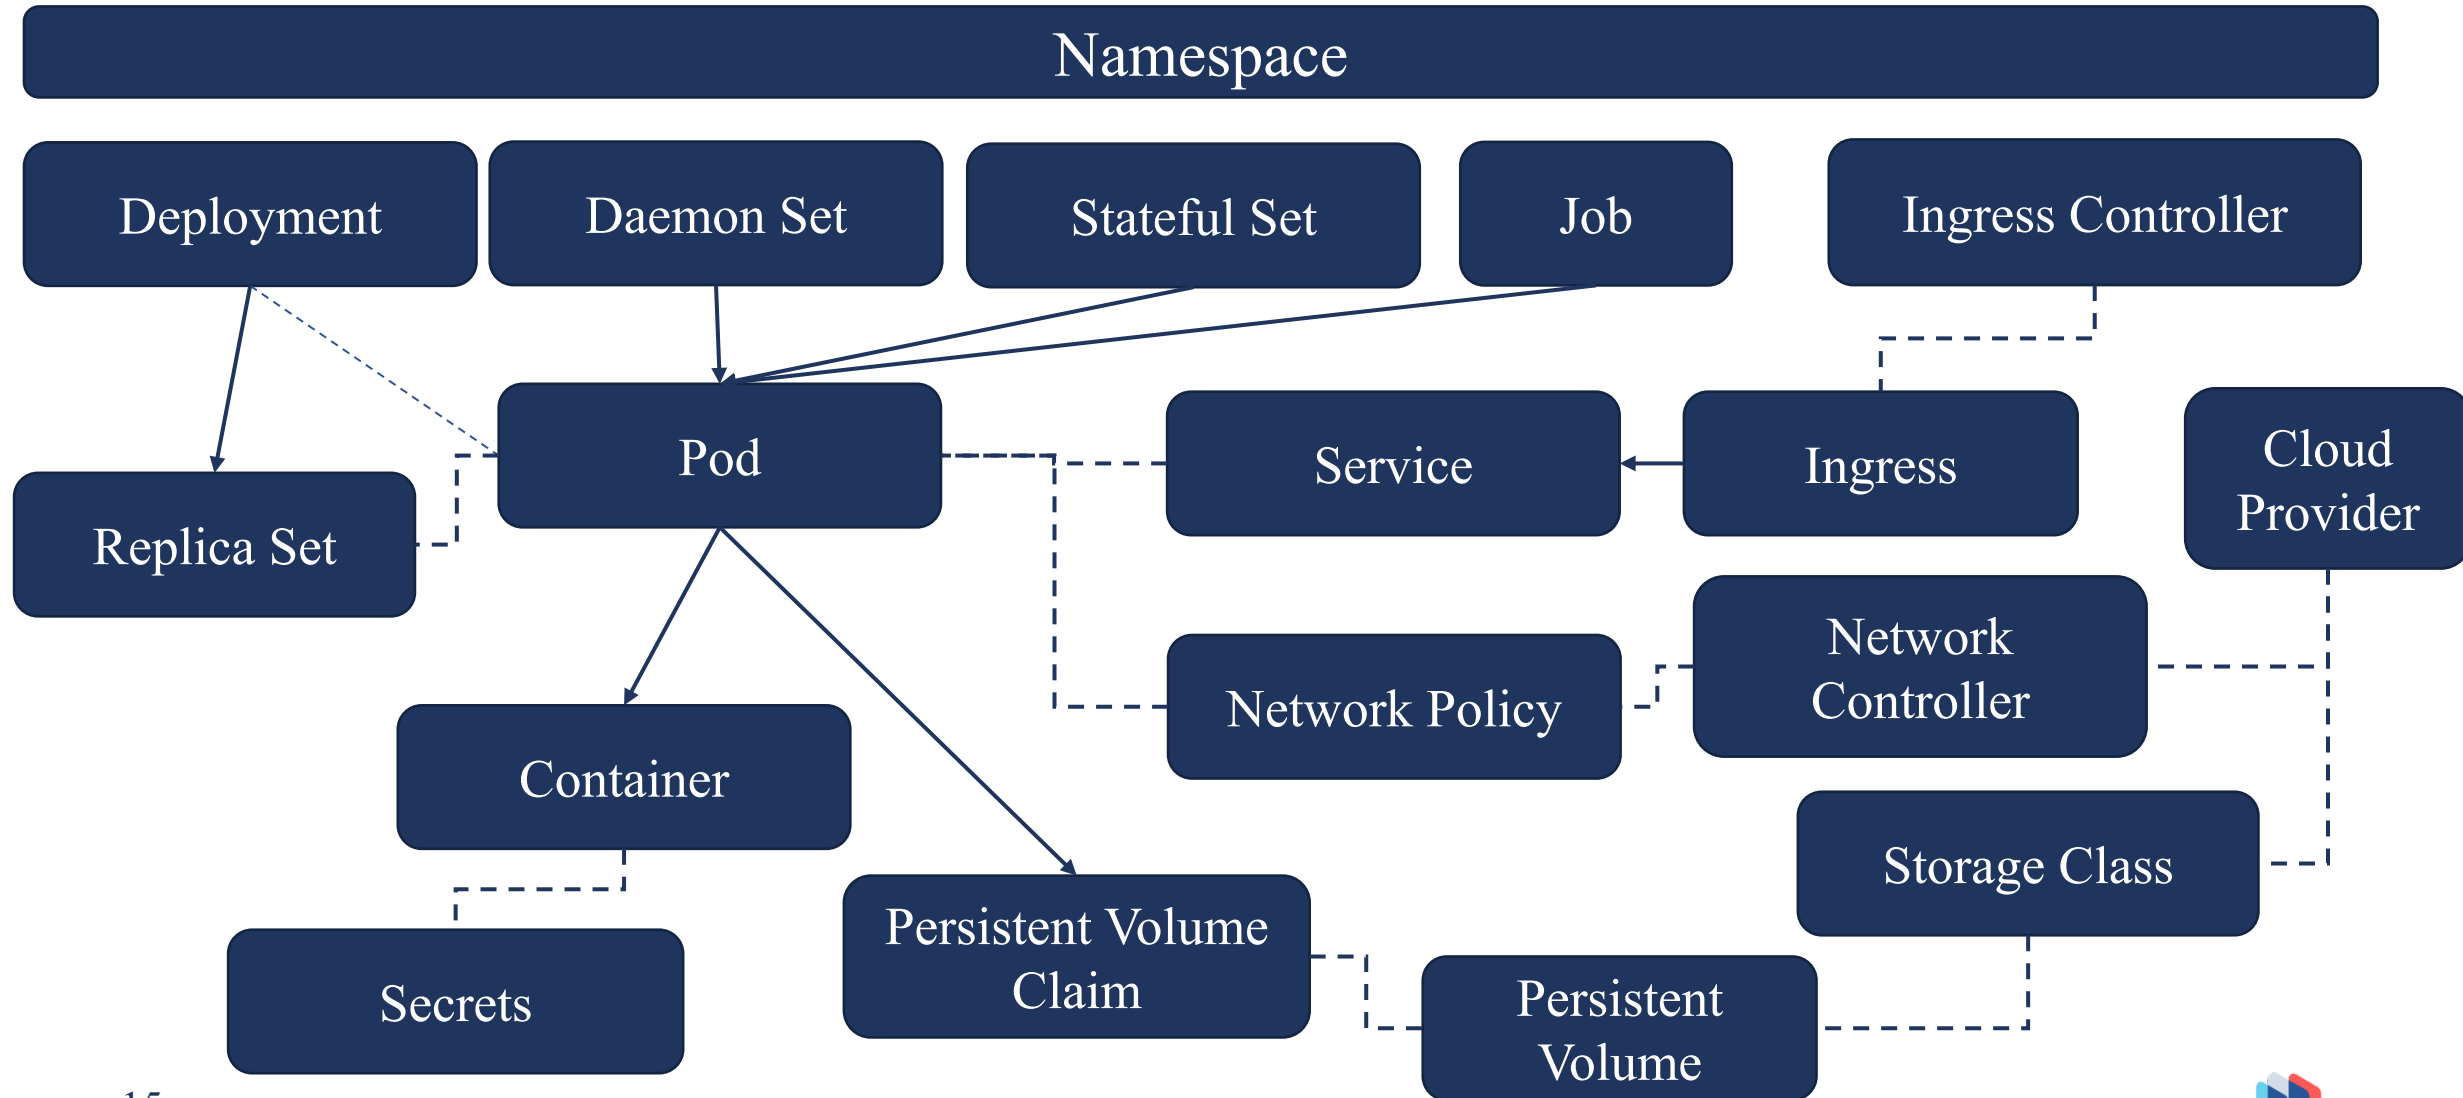
Setup your Azure Kubernetes Cluster



Kubernetes Application Concepts

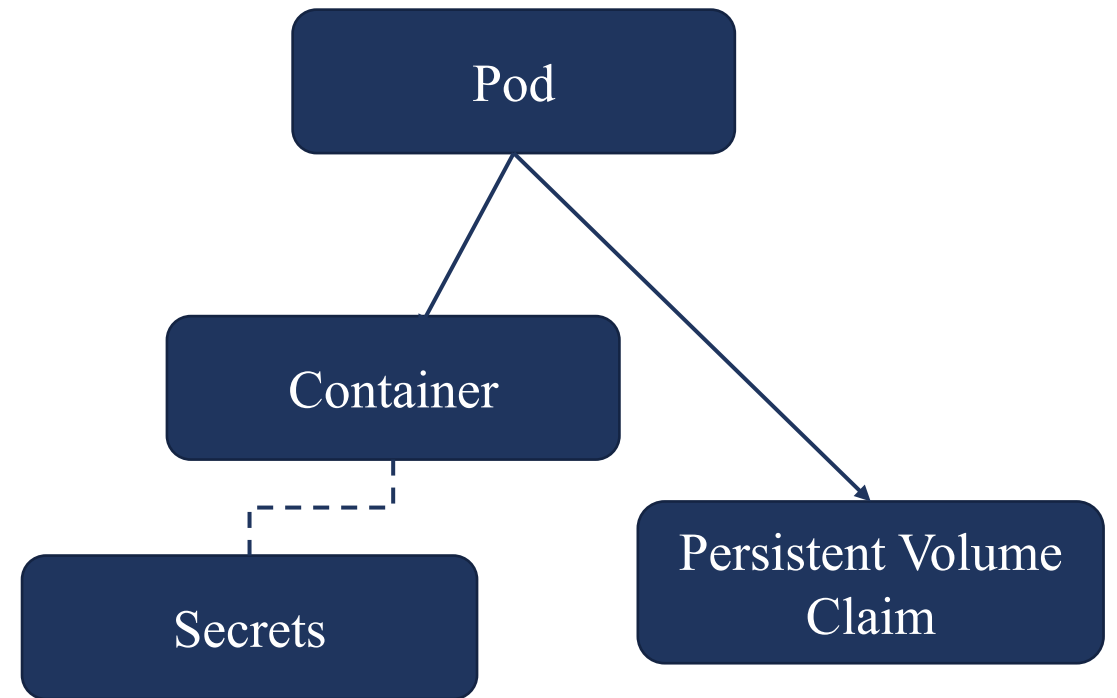


Dude, where's my app?



Pods

- Basic unit of application deployment
- Contains
 - One or more Containers
 - One or more PVCs
- Other constructs
 - nodeSelector
 - affinity
 - serviceAccountName
 - secrets
 - initContainers




Managing Pods

- Pods can be managed individually, but don't do this!
- Pods lifecycles are best managed using one of:
 - Deployments
 - StatefulSets
 - DaemonSets
- Less often used:
 - ReplicaSets (Deployments manage ReplicaSets)
 - Jobs (short-lived run-to-completion tasks)



Deployment

- Deployments automatically create (*and delete*) ReplicaSets
- Rollout: a new ReplicaSet is created and scaled up. The existing ReplicaSet is scaled down.
- Rollback: only impacts the Pod template. Can rollback to a specific revision ID.
- Rolling upgrade strategy tunables:
 - maxUnavailable
 - maxSurge



StatefulSet

- Pods with stable identities
 - names, network, storage
- Ordered creation, updates, scaling, and deletion
 - Pods are created, and named, in order from $\{0 \dots N-1\}$
- Use for clustered apps that use client-side identities
 - ZooKeeperAddresses: “zoo-1:2181, zoo-2:2181, zoo-3:2181”



DaemonSet

- Ensures that all Nodes run an instance of a Pod
- Useful for monitoring & security agents, log daemons, etc.
- A node selector can be used to target a subset of nodes

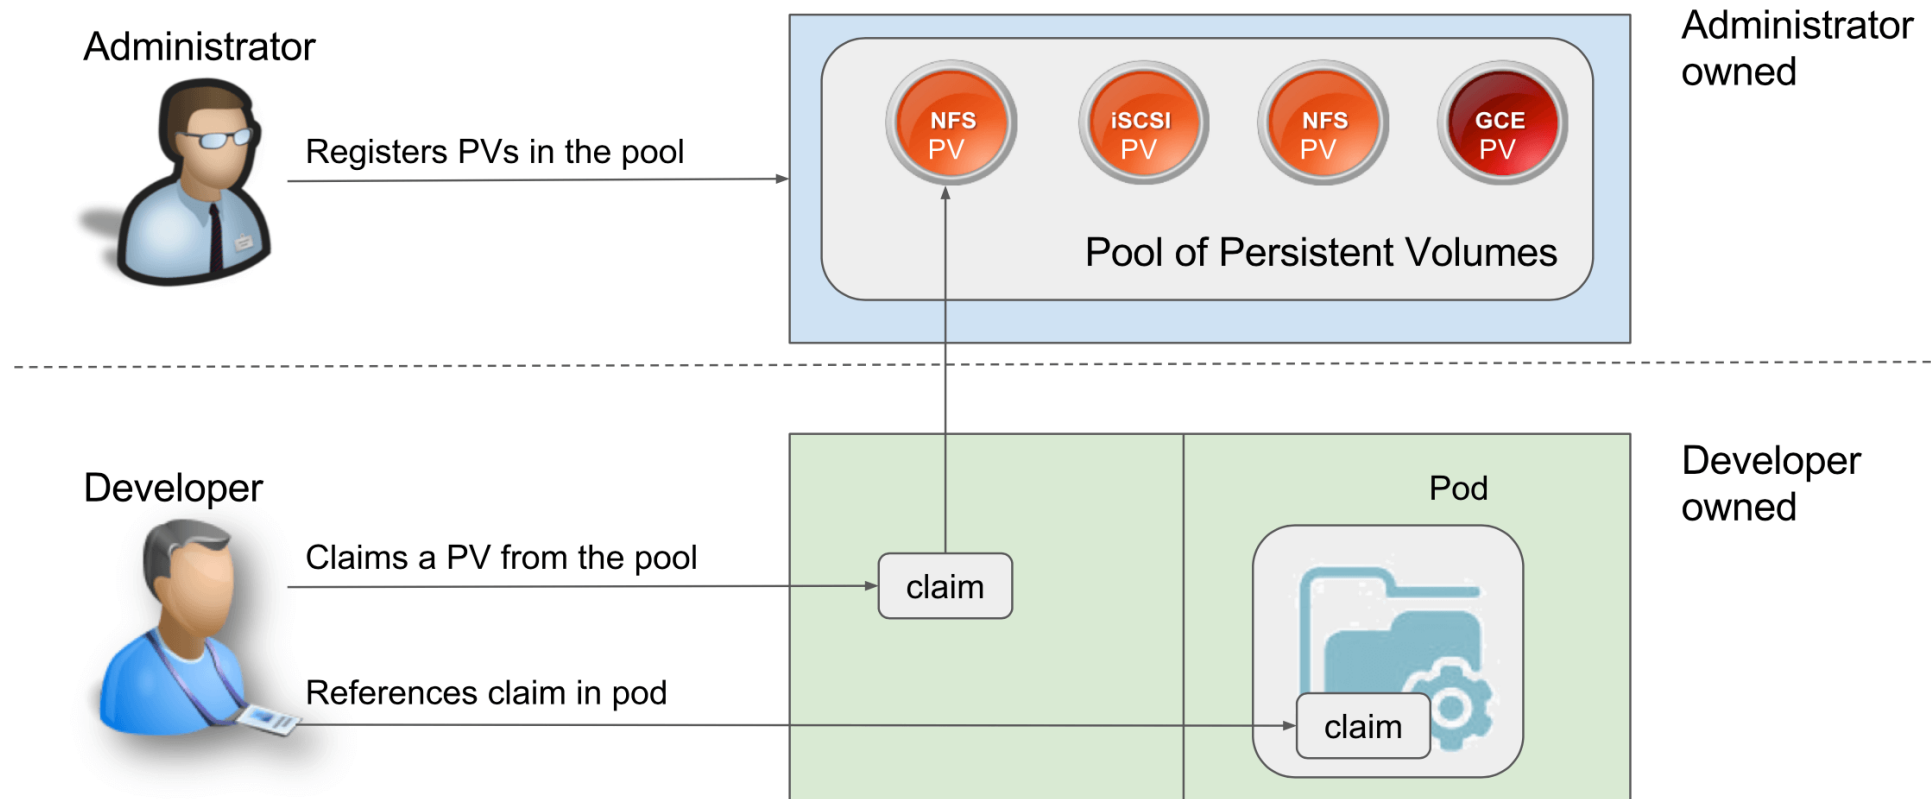


Networking your app

- **Service:** provides load-balancing. Addressed via IP (cluster IP) or a DNS name.
- **Network Policy:** manages routing rules across pods (east-west traffic.)
- **Ingress:** manages external routes to services (north-south traffic.) An Ingress Controller does the load-balancing. Ingress Resources specify the rules.



Configuring storage for your app



Source: Steve Watt, Red Hat

Summary

- Most apps will contain one or more services / tiers
- And each service will have:
 - Deployment → ReplicaSet → Pod → Container(s)
 - Service
 - Ingress *External facing services only*
 - Network Policy
 - Persistent Volume Claim(s)



Manage your Kubernetes Applications

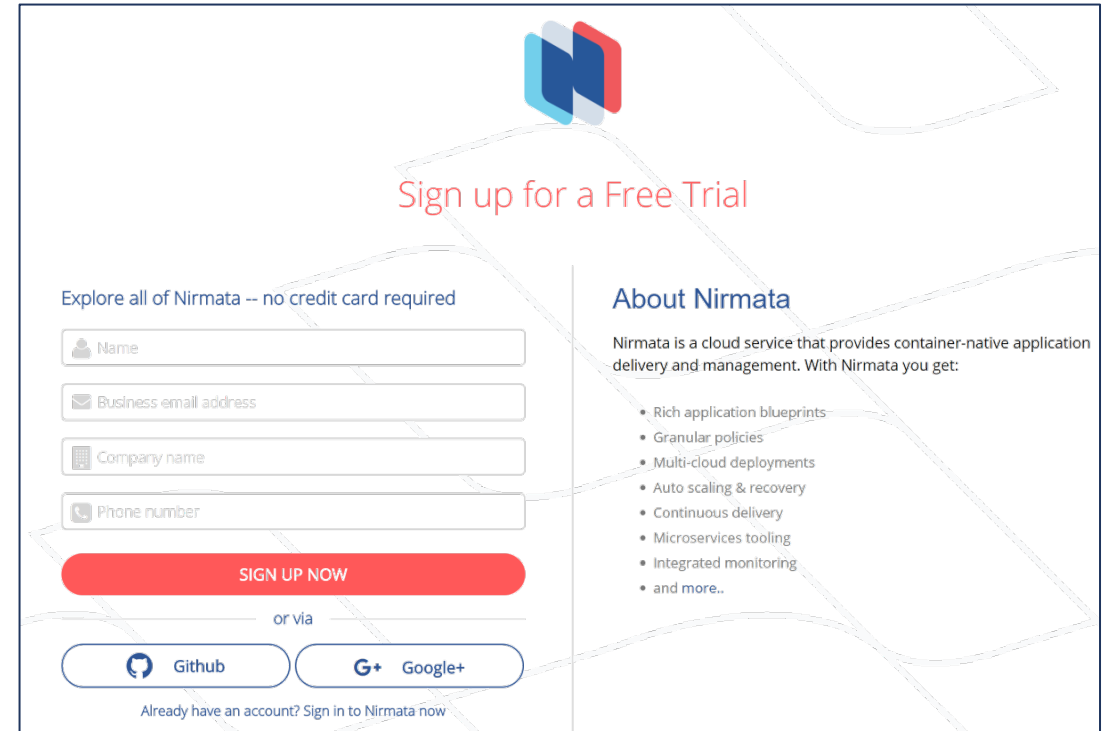


Signup for Nirmata

<https://try.nirmata.io>

Signup steps:

1. Enter your email
2. Nirmata will email you a secure access link
3. Click to setup your password and enter Nirmata!



The screenshot shows the Nirmata website's signup page. At the top center is the Nirmata logo, a stylized 'N' composed of blue and red geometric shapes. Below the logo, the text 'Sign up for a Free Trial' is displayed in red. The main content area is divided into two columns. The left column is titled 'Explore all of Nirmata -- no credit card required' and contains four input fields: 'Name' (with a person icon), 'Business email address' (with an envelope icon), 'Company name' (with a building icon), and 'Phone number' (with a phone icon). Below these fields is a prominent red button labeled 'SIGN UP NOW'. Underneath the button, it says 'or via' followed by two buttons for 'Github' and 'Google+'. At the bottom of this column, a link reads 'Already have an account? Sign in to Nirmata now'. The right column is titled 'About Nirmata' and contains a paragraph: 'Nirmata is a cloud service that provides container-native application delivery and management. With Nirmata you get:'. Below this paragraph is a bulleted list of features: 'Rich application blueprints', 'Granular policies', 'Multi-cloud deployments', 'Auto scaling & recovery', 'Continuous delivery', 'Microservices tooling', 'Integrated monitoring', and 'and more..'. The entire page has a light gray background with faint, abstract geometric lines.



Discover your Azure Cluster

1. Go to “Clusters → Add Cluster”
2. Select Kubernetes
3. Select Option 1, “Yes...” as we
4. Provide a name and select Other as provider
5. Download and apply the Nirmata Controller YAML



Deploy an application

1. Import the blueprint from Catalog: shopme-k8s
2. Deploy an Environment
 - Select Container Manager: Kubernetes
 - Select Type: <Cluster name>
 - Select Application: shopmek8s
 - Click on Finish



Thank-you!

<https://try.nirmata.com>

